

# Getting Started with KerasCV

# What is **Keras 3**?

- A full rewrite of of Keras
  - No technical debt
  - Much smaller codebase (43K LOC, ~3x smaller)
- Supports multi-backend (JAX, TensorFlow, PyTorch, and Numpy)
  - Numpy backend is inference only
- Drop in replacement for tf.keras when using TensorFlow backend
- Works seamlessly with KerasNLP and KerasCV



# Why use Keras 3?

- **One API** for every major framework (Helps avoid the fragmentation in the ecosystem)
- Seamlessly switch from one framework to another
  - Develop in the most intuitive one
  - Deploy in the fastest one
- Use `keras.ops` package to develop backend agnostic operations
  - Mimics the Numpy API (Same functions same arguments)
  - Includes extra functionalities for Neural Network (softmax, conv, cross\_entropy, etc.)
- Highly customizable



# One API

```
import os
os.environ["KERAS_BACKEND"] = "tensorflow"
import keras
keras.ops.numpy.arange(5)
>>> <tf.Tensor: shape=(5,), dtype=int32, numpy=array([0, 1, 2, 3, 4], dtype=int32)>
```

```
import os
os.environ["KERAS_BACKEND"] = "jax"
import keras
keras.ops.numpy.arange(5)
>>> Array([0, 1, 2, 3, 4], dtype=int32)
```

```
import os
os.environ["KERAS_BACKEND"] = "torch"
import keras
keras.ops.numpy.arange(5)
>>> tensor([0., 1., 2., 3., 4.]
```

# Advantages of Keras 3?

- Support for cross-framework data pipelines
- Pretrained models
- Progressive disclosure of complexity
- Introduces new stateless API for pure functional programming
- Distributed training as easy as non-distributed training



# Advantages of Keras 3?

- Support for cross-framework data pipelines
  - `tf.data.Dataset`
  - `torch.utils.data.DataLoader`
  - Numpy arrays
  - pandas dataframes
  - PyDatasets
- Pretrained models
- Progressive disclosure of complexity
- Introduces new stateless API for pure functional programming
- Distributed training as easy as non-distributed training

# Advantages of Keras 3?

- Support for cross-framework data pipelines
- Pretrained models
- **Progressive disclosure of complexity**
  - Start simple
  - Customize as per your needs
  - Go from Sequential/Functional to custom `train_step` to custom loops
- Introduces new stateless API for pure functional programming
- Distributed training as easy as non-distributed training

# KerasCV

KerasCV is a library of modular computer vision components that work natively with TensorFlow, JAX, or PyTorch, built on **Keras 3**.

```
pip install --upgrade keras-cv tensorflow
pip install --upgrade keras
```

Source: [https://keras.io/keras\\_cv/](https://keras.io/keras_cv/)



# Pre-trained models supported by KerasCV:

- **EfficientNetV2:** Various sizes of EfficientNet B-style architectures with different width and depth coefficients, boasting high accuracy on ImageNet.
- **MobileNetV3:** Large and small versions with hard-swish activation, optimized for mobile devices, pre-trained on ImageNet.
- **ResNet:** Both v1 and v2 versions with 50 layers, using batch normalization and ReLU activation, trained on ImageNet.
- **YOLOV8:** Different sizes of YOLOV8 backbones pre-trained on COCO for object detection tasks.

And more....

# Seamless Switching

```
import os
os.environ["KERAS_BACKEND"] = "tensorflow"
import keras
>>> Using TensorFlow backend
```

```
import os
os.environ["KERAS_BACKEND"] = "jax"
import keras
>>> Using JAX backend
```

```
import os
os.environ["KERAS_BACKEND"] = "torch"
import keras
>>> Using PyTorch backend
```

# Key Features:

- Offers models, layers, metrics, callbacks, and more that can be trained and serialized in any framework and re-used in another **without the need for costly migrations**.
- APIs **assist in common computer vision tasks** such as data augmentation, classification, object detection, segmentation, image generation, and more.
- Leverage **KerasCV** to quickly assemble **production-grade, state-of-the-art** training and inference pipelines for all of these common tasks.

Source: [https://keras.io/keras\\_cv/](https://keras.io/keras_cv/)

# From tensorflow import keras?

```
import tensorflow as tf
import tensorflow_datasets as tfds

from keras_cv.backend import keras ← Do this instead

import numpy as np

import keras_cv
from keras_cv import bounding_box
from keras_cv import visualization
from keras_cv.backend import ops
```

# Key Components:

- **Classification**
- **Object Detection Pipelines**
- **Image Classification and Augmentation**
- **Semantic Segmentation**
- **Image Generation with Stable Diffusion**

# Classification

Classification is the process of predicting a categorical label for a given input image. While classification is a relatively straightforward computer vision task. KerasCV provides APIs to construct commonly used components You can solve image classification problems at three levels of complexity:

- **Inference with a pre-trained classifier**

```
classifier = keras_cv.models.ImageClassifier.from_preset(  
    "efficientnetv2_b0_imagenet_classifier"  
)
```

Backbone model                      Task

- **Training a image classifier from scratch** (data augmentation, optimizer tuning, model, compile, fit)

```
backbone = keras_cv.models.EfficientNetV2B0Backbone() ← No pretrained weights  
model = keras.Sequential(  
    [  
        backbone,  
        keras.layers.GlobalMaxPooling2D(),  
        keras.layers.Dropout(rate=0.5),  
        keras.layers.Dense(101, activation="softmax"),  
    ]  
)
```

involved

- **Fine tuning pre-trained classifier**

When labeled images specific to our task are available, fine-tuning a custom classifier can improve performance. If we want to train a Cats vs Dogs Classifier, using explicitly labeled Cat vs Dog data should perform better than the generic classifier!

```
data, dataset_info = tfds.load("cats_vs_dogs", with_info=True, as_supervised=True)
train_steps_per_epoch = dataset_info.splits["train"].num_examples // BATCH_SIZE
train_dataset = data["train"]

num_classes = dataset_info.features["label"].num_classes

resizing = keras_cv.layers.Resizing(
    IMAGE_SIZE[0], IMAGE_SIZE[1], crop_to_aspect_ratio=True
)

def preprocess_inputs(image, label):
    image = tf.cast(image, tf.float32)
    # Staticly resize images as we only iterate the dataset once.
    return resizing(image), tf.one_hot(label, num_classes)

# Shuffle the dataset to increase diversity of batches.
# 10*BATCH_SIZE follows the assumption that bigger machines can handle bigger
# shuffle buffers.
train_dataset = train_dataset.shuffle(
    10 * BATCH_SIZE, reshuffle_each_iteration=True
).map(preprocess_inputs, num_parallel_calls=AUTOTUNE)
train_dataset = train_dataset.batch(BATCH_SIZE)
```

Pretrained backbones extract more information from our labeled examples by leveraging patterns extracted from potentially much larger datasets.

# Object Detection

**KerasCV** provides facilities to easily construct state-of-the-art object detection pipelines, offering pre-written data loaders and the ability to assemble production-grade data augmentation pipelines with **KerasCV** preprocessing layers.

```
pretrained_model = keras_cv.models.YOLOV8Detector.from_preset(  
    "yolo_v8_m_pascalvoc", bounding_box_format="xywh"  
)
```



```
bounding_boxes = {  
    "classes": [num_boxes],  
    "boxes": [num_boxes, 4]  
}
```

describes *exactly* what format the values in the "**boxes**" field of the label dictionary take in your pipeline.





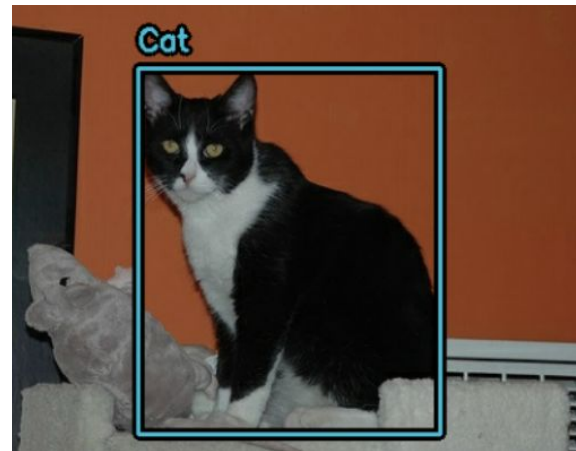
# Fine tuning pre-trained model

- Load Data

```
train_ds, ds_info = your_data_loader.load(  
    split='train', bounding_box_format='xywh', batch_size=8  
)
```

- Data Augmentation

```
augmenters = [  
    keras_cv.layers.RandomFlip(mode="horizontal", bounding_box_format="xywh"),  
    keras_cv.layers.JitteredResize(  
        target_size=(640, 640), scale_factor=(0.75, 1.3), bounding_box_format="xywh"  
    ),  
]  
  
def create_augmenter_fn(augmenters):  
    def augmenter_fn(inputs):  
        for augmenter in augmenters:  
            inputs = augmenter(inputs)  
        return inputs  
  
    return augmenter_fn
```



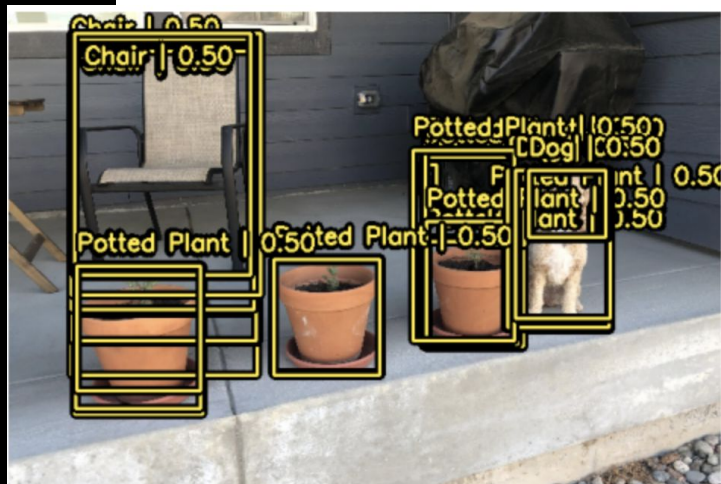
- Model

```
prediction_decoder = keras_cv.layers.NonMaxSuppression(
    bounding_box_format="xywh",
    from_logits=True,
    # Decrease the required threshold to make predictions get pruned out
    iou_threshold=0.2,
    # Tune confidence threshold for predictions to pass NMS
    confidence_threshold=0.7,
)

pretrained_model = keras_cv.models.YOLOV8Detector.from_preset(
    "yolo_v8_m_pascalvoc",
    bounding_box_format="xywh",
    prediction_decoder=prediction_decoder,
)

y_pred = pretrained_model.predict(image_batch)
visualization.plot_bounding_box_gallery(
    image_batch,
    value_range=(0, 255),
    rows=1,
    cols=1,
    y_pred=y_pred,
    scale=5,
    font_scale=0.7,
    bounding_box_format="xywh",
    class_mapping=class_mapping,
)
```

Non-max suppression is a traditional computing algorithm that solves the problem of a model detecting multiple boxes for the same object.



Check out this notebook by **Ian Stenbit**



# Data Augmentation

Data augmentation is a technique to make your model robust to changes in input data such as lighting, cropping, and orientation. KerasCV includes some of the most useful augmentations in the `keras_cv.layers` API.

**KerasCV** supports training powerful image classifiers and includes a suite of preprocessing layers implementing common data augmentation techniques such as CutMix, MixUp, and RandAugment.

# Common Augmentations

- RandomFlip()



- RandomCropAndResize()



- RandomCutout()

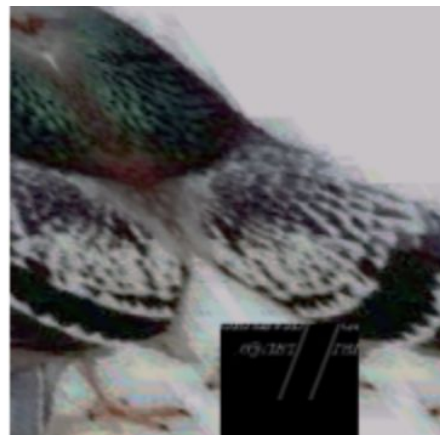
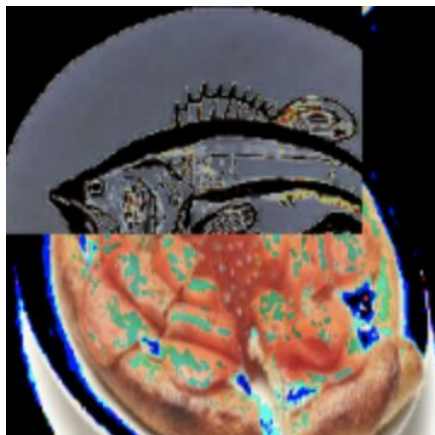


# CutMix

Instead of replacing the cut-out areas with black pixels, **CutMix** replaces these regions with regions of other images sampled from within your training set! Following this replacement, the image's classification label is updated to be a blend of the original and mixed image's class label.

```
cut_mix = keras_cv.layers.CutMix()
# CutMix needs to modify both images and labels
inputs = {"images": image_batch, "labels": label_batch}

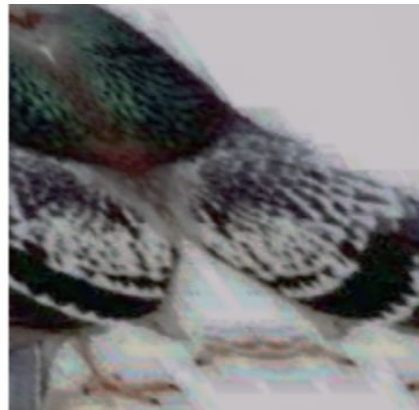
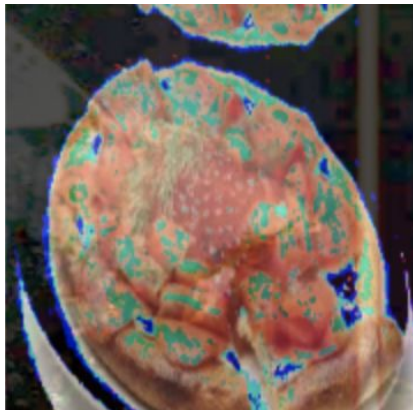
keras_cv.visualization.plot_image_gallery(
    cut_mix(inputs)["images"],
    rows=3,
    cols=3,
    value_range=(0, 255),
    show=True,
)
```



# MixUp

`MixUp()` works by sampling two images from a batch, then proceeding to literally blend together their pixel intensities as well as their classification labels.

```
mix_up = keras_cv.layers.MixUp()  
# MixUp needs to modify both images and labels  
inputs = {"images": image_batch, "labels": label_batch}  
  
keras_cv.visualization.plot_image_gallery(  
    mix_up(inputs)["images"],  
    rows=3,  
    cols=3,  
    value_range=(0, 255),  
    show=True,  
)
```



# Segmentation

Semantic segmentation is a type of computer vision task that involves assigning a class label such as person, bike, or background to each individual pixel of an image, effectively dividing the image into regions that correspond to different object classes or categories.

DeepLabv3+ developed by Google for semantic segmentation, combining advanced techniques for accurate and detailed segmentation results.

## Semantic segmentation with a pretrained DeepLabv3+ model

```
model = keras_cv.models.DeepLabV3Plus.from_preset(  
    "deeplab_v3_plus_resnet50_pascalvoc",  
    num_classes=21,  
    input_shape=[512, 512, 3],  
)
```



Image



Ground Truth



Prediction



Image



Ground Truth



Prediction



Image



Ground Truth



Prediction



Image



Ground Truth



Prediction



# Image Generation with Stable Diffusion

Stable Diffusion is a powerful image generation model that can be used, among other things, to generate pictures according to a short text description (called a "prompt"). To generate novel images based on a text prompt using the KerasCV implementation of [stability.ai](#)'s text-to-image model, [Stable Diffusion](#).

```
model = keras_cv.models.StableDiffusion(  
    img_width=512, img_height=512, jit_compile=False  
)
```

```
images = model.text_to_image("photograph of an grinch on christmas", batch_size=3)  
import matplotlib.pyplot as plt  
  
def plot_images(images):  
    plt.figure(figsize=(20, 20))  
    for i in range(len(images)):  
        ax = plt.subplot(1, len(images), i + 1)  
        plt.imshow(images[i])  
        plt.axis("off")
```



# Comparison with 🙌 Diffusers

Both implementations were tasked to generate 3 images with a step count of 50 for each image with a Tesla T4 GPU.

GPU	Model	Runtime
Tesla T4	KerasCV (Warm Start)	<b>28.97s</b>
Tesla T4	diffusers (Warm Start)	41.33s
Tesla V100	KerasCV (Warm Start)	<b>12.45</b>
Tesla V100	diffusers (Warm Start)	12.72

However, the results are still not comparable for cold start. Cold-start execution time includes the one-time cost of model creation and compilation, and is therefore negligible in a production environment (where you would reuse the same model instance many times).

Note that the StableDiffusion API, as well as the APIs of the sub-components of StableDiffusion (e.g. ImageEncoder, DiffusionModel) should be considered unstable at this point. We do not guarantee backwards compatibility for future changes to these APIs.

# Resources

- Keras\_CV Github: <https://github.com/keras-team/keras-cv>
- Classification: [https://keras.io/guides/keras\\_cv/classification\\_with\\_keras\\_cv/](https://keras.io/guides/keras_cv/classification_with_keras_cv/)
- Object Detection: [https://keras.io/guides/keras\\_cv/object\\_detection\\_keras\\_cv/](https://keras.io/guides/keras_cv/object_detection_keras_cv/)
- Image Augmentation: [https://keras.io/guides/keras\\_cv/custom\\_image\\_augmentations/](https://keras.io/guides/keras_cv/custom_image_augmentations/)
- Segmentation: [https://keras.io/guides/keras\\_cv/semantic\\_segmentation\\_deeplab\\_v3\\_plus/](https://keras.io/guides/keras_cv/semantic_segmentation_deeplab_v3_plus/)
- Image Generation with Stable diffusion:  
[https://keras.io/guides/keras\\_cv/generate\\_images\\_with\\_stable\\_diffusion/](https://keras.io/guides/keras_cv/generate_images_with_stable_diffusion/)

THANK YOU